

Parallel Barnes-Hut algorithm to deal with N-Body problem

Master Erasmus EPFL Student

ATTACHE William

`attachew974@gmail.com`

Linkedin

Département de Mathématiques - MATH454
Ecole Polytechnique Fédérale de Lausanne

Spring Semester 2018

Outline

Introduction

Architecture of the application

Difficulties & improvements

Scaling

Conclusion



Introduction

- N-Body problem : *predict the individual motions of a group of objects interacting with each other under gravitational forces*
- Parallelize : Sequential $O(n^2)$ v.s. Barnes-Hut $O(n * \ln(n))$
- Work load balancing
- Communication between processors : MPI
- Local Essential Tree (LET)

Parallelization choices

- ① gprof and code profiling (20000 particules, 1000 time steps)

Method	% of total time	Number of calls
<i>insertNode</i>	78,62	5483872
<i>computeForce</i>	9,00	5000000
<i>gForce</i>	6,77	80041997
<i>dist</i>	4,91	108258504
<i>getSubtree</i>	3,38	80084009
<i>buildTree</i>	0,75	1000

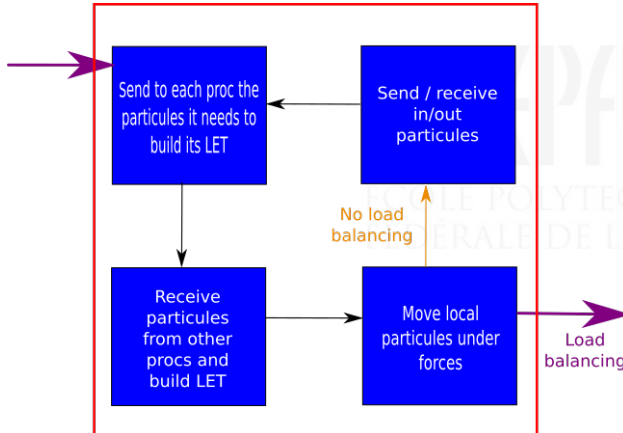
- ② MPI \ OpenMP

Communications

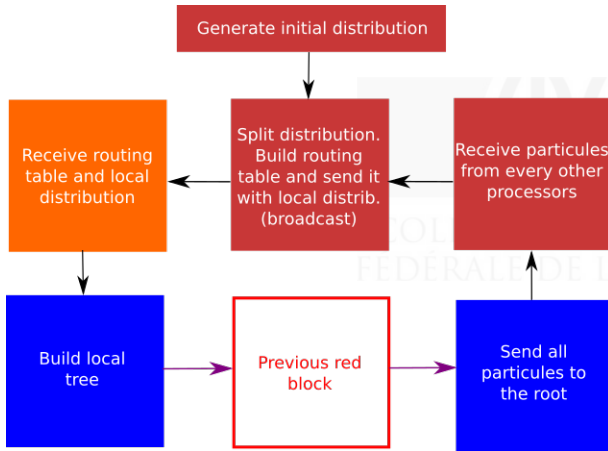
- ① Routing table (*MPI_Broadcast*) & local distribution
- ② Load balancing (*MPI_Send*\ *MPI_Recv*)

Global architecture

For a processor with its local distribution.



Global architecture



Split the work load among processors

```
#areas = 1 ; #leftprocs = p-1
```




Split the work load among processors

```
#areas = 2 ; #leftprocs = p-2
```




Split the work load among processors

```
#areas = 3 ; #leftprocs = p-3
```




Split the work load among processors

```
#areas = 4 ; #leftprocs = p-4
```




Split the work load among processors

`#areas = 5 ; #leftprocs = p-5`




Split the work load among processors

```
#areas = 6 ; #leftprocs = p-6
```




Split the work load among processors

`#areas = 7 ; #leftprocs = p-7`




Split the work load among processors

`#areas = 8 ; #leftprocs = n-8`



Split the work load among processors

`#areas = 9 ; #leftprocs = n-9`



Main differences with Berkeley lecture notes

- ① not only for $p = 2^m$ processors
- ② do not need to perform a round-robin to send particules :
routing table access in $O(1)$
 - ▷ **main drawback** (reasonable) memory consumption to store such a table for every processor

Difficulties encountered

- ① "learn" C++ and so deal with memory \ scope : `addChilds`
- ② lots of *absent-mindedness* (move particules, reshape boxes)
- ③ particules received but coordinates changed
→ define a `MPI_Data_type` *particules* as a solution
- ④ *synchronization point* : impossible to run with more than 2500 particules
- ⑤ write at the right place in I \ O file

Improvements considered

- 1 Recompute the center of mass when a particule moves
- 2 Run simulations with heterogeneous distributions
- 3 Manage particule collisions and particules that go out
- 4 Add genericity to implement quick sort : define *comparison operators*.
- 5 Use more *auxiliary functions* in main, like fill arrays with addresses (build rooting table, back sending for load balancing) ; use more classes to structure the code.
- 6 Do not systematically create four childs when building the tree to save (a little) of memory

Manage particle collisions and particles that go out

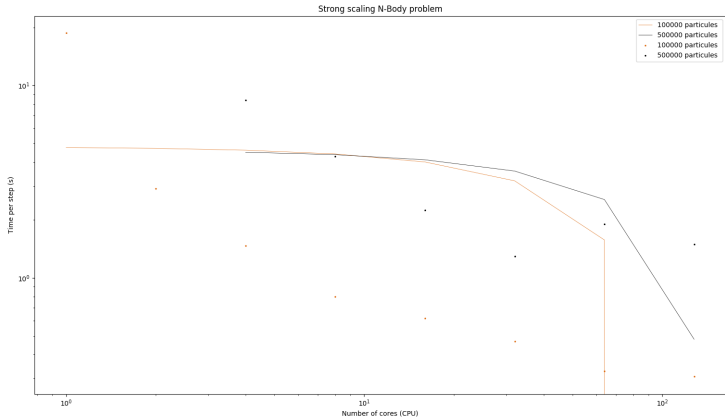
- ▷ instead of going out, a particle could :
 - bounce on the wall
 - (erroneously) be considered as a point that goes away with the direction and the velocity it had when leaving the area
- ▷ collisions : for any particule, consider all particules which could have ramed it → add technical difficulties and computations

Strong scaling

- ▷ behaviour of the application by fixing the total size of the problem and increasing the number of processors

$$\text{SPEEDUP} : S_p(n) = \frac{t_p(n)}{t_1(n)}$$

Strong scaling




Weak scaling

- ▷ behaviour of the application by fixing the total size of the problem and increasing the number of processors

$$\text{PARALLEL EFFICIENCY} : E_p = \frac{S_p(n)}{p}$$

Weak scaling



Conclusion

- Encouraging scaling results
- But lots a work left to get a proper application
- Deal with particules collisions could obstruct scaling
- Maybe CUDA could offer better scaling results